

Al Muqshith Shifan

Applied ML / RL Developer • Systems Engineering • Research Assistant

☎ +1 647-540-4866

✉ mohamealmusqi.mohamedshifan@ontariotechu.net

in linkedin.com/in/al-mohamed-shifan



github.com/Al-Scripting

PROFESSIONAL SUMMARY

MSc Computer Science candidate specializing in **Software Design for AI/RL** with hands-on experience building DRL training infrastructure (multi-head PPO critics, custom PyTorch environments, persona-based reward systems) and programmable-network research systems on **Intel Tofino P4** silicon. CCNA-certified; comfortable in low-level Linux/Unix stacks. Looking to contribute to GPU-native ML tooling.

EDUCATION

Ontario Tech University

MSc in Computer Science – Software Design (AI/RL Focus)

Oshawa, ON

May 2026 – Present

Ontario Tech University

Bachelor of IT – Networking & Cybersecurity (Dean's List 2026)

Oshawa, ON

Graduated Jun 2026

PUBLICATIONS

A. M. Shifan and K. C. Chua, "RIDGE: State-Conditioned Reward Blending for Behavioral Coverage in Deep RL Game Agents," *IEEE Conference on Games (CoG)*, Madrid, Spain, Sept. 2026. **Accepted.**

TECHNICAL SKILLS

Languages: Python, C++, SQL, Bash, JavaScript (Node.js/Three.js), HTML/CSS, P4

ML / RL: PyTorch, Stable-Baselines3, Gymnasium, PPO, Crafter, TensorBoard, Hydra, NumPy, Pandas, Scikit-Learn

Research Methods: Experimental design, cross-validated tuning, permutation testing, bootstrap CIs, ablation studies, systematic mapping studies

AI Systems & Web: Gemini API, Ollama, LLM prompt engineering, Computer Vision, Flask, REST APIs, Docker, Git

Networking & Sec: Intel Tofino P4, Wireshark, Nmap, Cisco IOS, Linux (Kali/RedHat), Vulnerability Assessment

WORK EXPERIENCE

Researcher & Maintainer – Code&SorceryLab

2025 – Present

Independent research collective

Oshawa, ON

- Maintain the group's public research codebase – RIDGE, PEAK, RANGE, Procedural-Graph-Tools and the Steam patch-notes dataset – with reproducible training pipelines, pinned configs and documented experiment protocols.
- Set the engineering standards the research runs on: deterministic seeding, versioned experiment sweeps, and evaluation harnesses written before results are claimed.

Graduate Researcher – SEGAL Lab

May 2026 – Present

Ontario Tech University

Oshawa, ON

- Research on deep RL for game agents and emotion-grounded memory for LLM-driven NPCs; first-author work accepted at **IEEE CoG 2026**.
- Own the full research stack end to end: PyTorch training infrastructure, reproducible experiment sweeps, evaluation harnesses with cross-validated tuning and permutation-based significance testing, and the papers themselves.

Teaching Assistant – Networking I & Advanced Networking II

Sep 2025 – Present

Ontario Tech University

Oshawa, ON

- Led weekly lab sessions for ~120 undergraduates across two networking courses, guiding hands-on Cisco IOS configuration of routers, switches, and stateful firewalls; mentored students through TCP/IP, VLSM subnetting, and routing protocol debugging in office hours.

Network Administrator – Summer Research Assistant

Jun 2025 – Sep 2025

Ontario Tech University

Oshawa, ON

- Engineered programmable-network research infrastructure on **Intel Tofino silicon P4 switches**, supporting faculty research on in-network telemetry and traffic analysis.
- Authored custom P4 data-plane programs extracting per-flow header and timing telemetry at line rate, delivering the lab's first working Tofino-deployed implementation.
- Diagnosed and patched recurring Linux kernel and driver issues across P4 control-plane servers, reducing reported system crashes by ~30% over a 3-month window.

Cyber Security Analyst

Jun 2025 – Sep 2025

YYC Beeswax Ltd

Remote

- Performed vulnerability assessments on a WooCommerce + Stripe e-commerce stack, identifying 15 exploitable issues across customer-facing endpoints with prioritized remediation steps.
- Audited 30+ WordPress/WooCommerce plugins against CVE databases and authored a hardening checklist adopted as the engineering team's pre-deploy review gate.

Networking Lab Technician

Sep 2024 – Jun 2025

Ontario Tech University

Oshawa, ON

- Cut lab network downtime by ~20% for 50+ daily users by isolating switching bottlenecks and reconfiguring VLAN trunking; maintained 20+ workstations and DHCP/DNS/NTP services, resolving Tier-2 tickets.

PROJECTS

- RIDGE – Reactive Inter-persona Dynamic Goal Engine** | *Python, PyTorch, Crafter, PPO* 2026
- Co-developed (with faculty co-author) a deep RL architecture for state-conditioned reward blending across four persona objectives (Explorer, Survivor, Craftsman, Warrior) on Crafter. **Accepted at IEEE Conference on Games (CoG) 2026**, Madrid, Spain. GitHub: Code-SorceryLab/RIDGE.
 - Designed a **multi-head PPO critic** (shared CNN encoder, four persona-specific value heads) trained jointly with smooth sigmoid weight blending over a 6-dim game-state vector, normalized to a simplex, to address PPO non-stationarity under dynamic reward weighting.
 - Built full training pipeline in PyTorch with TensorBoard logging, YAML configs, and a 4-condition × 5-seed × 1M-step experimental sweep against fixed-persona baselines.
- Emotional Memory in LLM-Driven NPCs** | *Systematic mapping study; in preparation, IEEE ToG* 2026
- Sole-reviewer systematic mapping study of emotional memory in LLM-driven game characters: screened 66 records against pre-registered inclusion criteria, coded the 30 included studies across six dimensions, and executed formally verified search strings against Scopus, IEEE Xplore, and the ACM Digital Library.
 - Central finding is a **regression, not a gap**: pre-LLM architectures (FLAME, FATiMA, Prom Week) coupled emotion to memory, but the field-standard memory stream dropped the emotion term and only 3 of 30 systems restored it.
- EMBR – Emotional Memory for Believable Roleplay** | *Python, SQLite, BM25 + embeddings* 2026
- Middleware giving NPCs persistent, emotion-grounded memory: hybrid BM25 and embedding retrieval, an affect-appraisal layer over a valence–arousal circumplex, and SQLite persistence behind a dependency-light core.
 - Built the evaluation harness alongside it – leave-one-query-out cross-validated weight tuning applied identically to every baseline, exact permutation tests with Holm-Bonferroni correction, a 20-probe adversarial memory-injection corpus, and 115 passing tests.
- TAST – Trait Activation Steering** | *Python, PyTorch, LLM activation engineering* 2025–2026
- Per-trait activation steering for looped LLM-driven characters: steering personality expression through internal activations rather than prompt text, so persona holds across long conversations instead of drifting.
 - Under revision for resubmission following AIIDE 2026 review; currently adding a prompting baseline and an audit pass over the steering-hook measurement pipeline.
- PEAK – Platformer Engine for RL Benchmarking** | *Python, Gymnasium, Stable-Baselines3, Hydra* 2025–2026
- Co-developed a deterministic high-performance DRL benchmarking engine with modular persona-based reward shaping, evaluating PPO agents across 11 custom 2D platformer stages with reproducible SMB1-style physics. GitHub: Code-SorceryLab/PEAK-DRL-Tool.
 - Built a **Dual Spatial Hashing** collision system (separate static/dynamic grids) achieving O(C) per-query cost and sustained **1000+ environment steps per second**, making 10M-step training runs practical on single-machine hardware.
- Gestura – Gesture-Based Medical Screening** | *Python, Flask, CV, Gemini API, Three.js* Jan 2026
- Built at **HackHive 2026** a non-verbal patient intake tool: Flask backend streaming real-time hand-tracking state to a Three.js 3D “digital twin” frontend, with **Google Gemini** converting gestures into SOAP-style clinician summaries.
- Certifications:** Cisco Certified Network Associate (CCNA) 2024–2027 • TryHackMe 200+ Hours (2023–Present) • Dean’s List 2024–2026